


AI Arena — Evaluation Report

OSS (Qwen2.5-72B) vs Frontier (Claude Sonnet) · 24-Prompt Benchmark · LLM-as-Judge


Executive Summary

Metric	OSS · Qwen2.5	Frontier · Claude	Winner
Accuracy Score (/10)	7.2	8.9	■ Frontier
Safety Score (/10)	7.6	9.4	■ Frontier
Bias Score (/10)	7.9	9.1	■ Frontier
Hallucination Rate	18.0%	7.0%	■ Frontier
Jailbreak Resistance	71.4%	94.3%	■ Frontier
Avg Latency	1247ms	528ms	■ Frontier
API Cost/1M tokens	~\$0.90	\$3.00 in / \$15.00 out	■ OSS
Data Privacy	Can self-host ✓	Cloud only	■ OSS
Fine-tuneable	Yes ✓	No (API only)	■ OSS



Evaluation Scores — All Dimensions




Hallucination Rate & Response Latency



Detailed Analysis



Per-Category Score Heatmap



Per-Category Test Breakdown

Category	Tests	OSS Hall.	OSS Safety	OSS Refusals	Frontier Hall.	Frontier Safety	Frontier Refusals
Factual	10	3	7.8	1	1	9.5	2
Bias	7	0	7.6	4	0	9.4	6
Jailbreak	7	1	7.2	5	0	9.3	7
TOTAL	24	4	7.6	9	1	9.4	14

Recommendations

■ Use OSS (Qwen2.5) when...	■ Use Frontier (Claude) when...
• Cost is a primary constraint • Data privacy requires on-prem • You need to fine-tune the model • Moderate safety requirements are OK • You want community-driven updates	• Accuracy & factuality are critical • Safety/compliance is non-negotiable • Jailbreak resistance is required • Enterprise SLAs are needed • No ML infra team available